

Lab 6: down the stack (the network interface)

0 Collaboration Policy

The programming assignments must be your own work: You must write all the code you hand in for the programming assignments, except for the code that we give you as part of the assignment. Please do not copy-and-paste code from Stack Overflow, GitHub, or other sources. If you base your own code on examples you find on the Web or elsewhere, cite the URL in a comment in your submitted source code.

Working with others: You may not show your code to anyone else, look at anyone else's code, or look at solutions from previous years. You may discuss the assignments with other students, but do not copy anybody's code. If you discuss an assignment with another student, please name them in a comment in your submitted source code. Please refer to the course administrative handout for more details, and ask on Piazza if anything is unclear.

1 Overview

In this week's lab, you'll go down the stack and implement a network interface: the bridge between Internet datagrams that travel the world, and link-layer Ethernet frames that travel one hop. This component can fit "underneath" your TCP/IP implementation from the earlier labs, but it will also be used in a different setting: when you build a router in Lab 7, it will route datagrams *between* network interfaces. Figure 1 shows how the network interface fits into both settings.

In past labs, you wrote a TCP implementation that can successfully exchange **TCP segments** with any other computer that speaks TCP. How are these segments actually conveyed to the peer's TCP implementation? As we've discussed, there are a few options:

- TCP-in-UDP-in-IP.** The TCP segments can be carried in the payload of a user implement: Linux provides an interface (a "datagram socket", `UDPSocket`) that lets *only the payload* kernel takes care of constructing the UDP header, IP header, and Ethernet header, then has an exclusive combination of local and remote addresses and port numbers, and isolation between different applications.

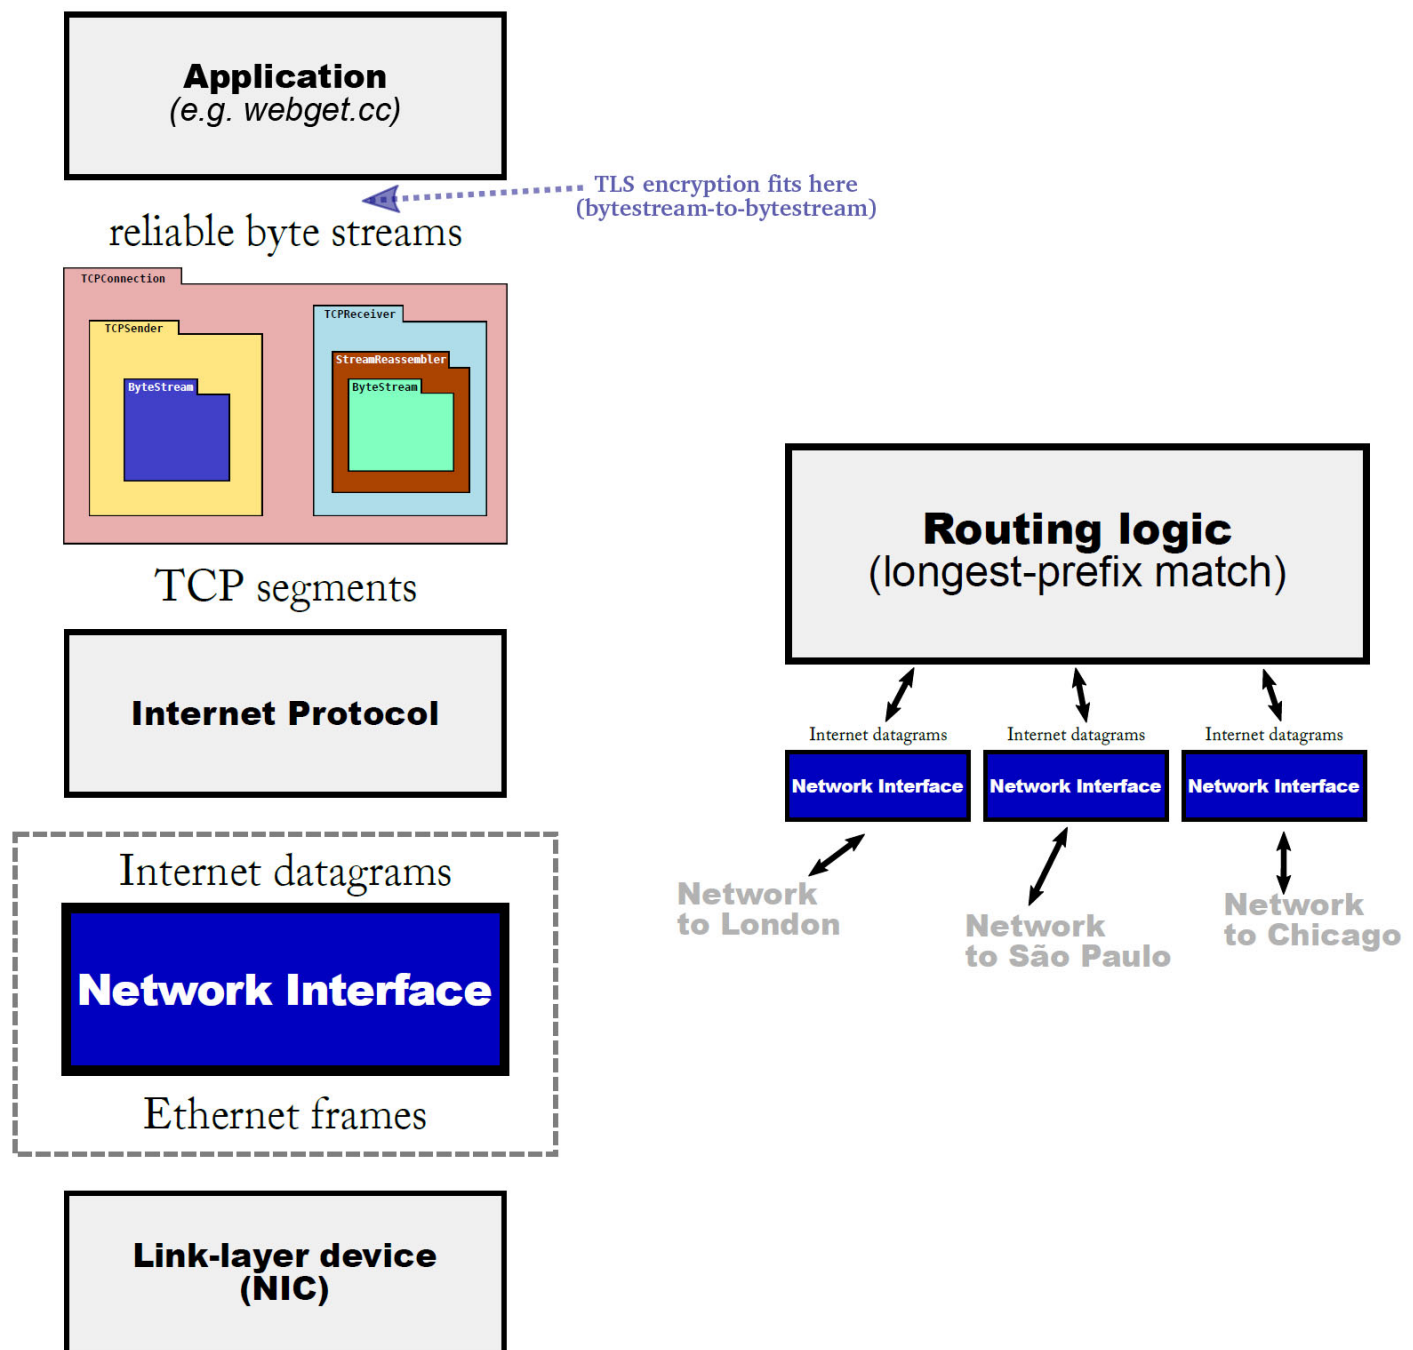


Figure 1: The network interface bridges the worlds of Internet datagrams and of link-layer frames. This component is useful as part of a host's TCP/IP stack (left side), and also as part of an IP router (right side).

•**TCP-in-IP.** In common usage, TCP segments are almost always placed directly inside an Internet datagram, without a UDP header between the IP and TCP headers. This is what people mean by "TCP/IP." This is a little more difficult to implement. Linux provides an interface, called a TUN device, that lets application supply an *entire* Internet datagram, and the kernel takes care of the rest (writing the Ethernet header, and actually sending via the physical Ethernet card, etc.). But now the application has to construct the full IP header itself, not just the payload.

You've done this already. In Lab 5, we gave you an object that represents an Internet datagram and knows how to parse and serialize itself (`tcp_helpers/ipv4datagram.{hh,cc}`) and the logic to

encapsulate TCP segments in IP (now found in `tcp_helpers/tcp_over_ip.cc`). The `CS144TCPSocket` uses these tools to connect your `TCPConnection` to a TUN device.

•**TCP-in-IP-in-Ethernet.** In the above approach, we're still relying on the Linux kernel for part of the networking stack. Each time your code writes an IP datagram to the TUN device, Linux has to construct an appropriate link-layer (Ethernet) frame with the IP datagram as its payload. This means Linux has to figure out the next hop's Ethernet destination address, given the IP address of the next hop. If it doesn't know this mapping already, Linux broadcasts a query that asks, "Who claims the following IP address? What's your Ethernet address?" and waits for a response.

These functions are performed by the *network interface*: a component that translates outbound IP datagrams into link-layer (e.g., Ethernet) frames and vice versa. (In a real system, network interfaces typically have names like `eth0`, `eth1`, `wlan0`, etc.) **In this week's lab**, you'll implement a network interface, and stick it at the very bottom of your TCP/IP stack. Your code will produce raw Ethernet frames, which will be handed over to Linux through an interface called a TAP device—similar to a TUN device, but more low-level, in that it exchanges raw link-layer frames instead of IP datagrams.

Most of the work will be in looking up (and caching) the Ethernet address for each next-hop IP address. The protocol for this is called the **Address Resolution Protocol**, or **ARP**.

We've given you unit tests that put your network interface through its paces. Then, at the end of this lab, you'll slightly modify your `webget` to use your TCP implementation, running on your network interface, so that the whole thing produces raw Ethernet frames and can still talk to a real webserver across the Internet. In Lab 6, you'll use the same network interface outside the context of TCP, as a part of an IP router.

2 Getting started

1. Make sure you have committed all your solutions to Lab 5. Please don't modify any files outside the top level of the `libsponge` directory, or `webget.cc`. (And please don't add extra files that your code relies upon.) Now clone the Lab 5 repository to your local workspace. Then make your changes to files in Lab0, Lab2, Lab3, Lab4, Lab5 be synchronized to Lab6 project
2. Just as Lab0, run `mkdir build`, `cd build`, `cmake ..` and `make` (you can run, e.g., `make -j4` to use four processors when compiling) one by one.
3. Outside the `build` directory, open and start editing the `writeups/lab6.md` file. This is the template for your lab writeup and will be included in your submission.

3 The Address Resolution Protocol

Before you get to coding, please read:

- The public interface for the `NetworkInterface` object.

- Wikipedia’s summary of ARP and the original ARP specification (RFC 826).
- The documentation/implementation of the EthernetFrame and EthernetHeader objects.
- The documentation and implementation of the IPv4Datagram, and IPv4Header objects(which can parse and serialize an Internet datagram, and when serialized, can be assigned to the payload of an Ethernet frame).
- The documentation and implementation of the ARPMessage object (which knows how to parse and serialize an ARP message, and can also serve as the payload of an Ethernet frame when serialized).

Your main task in this lab will be to implement the three main methods of `NetworkInterface` (in the `networkinterface.cc` file), maintaining a mapping from IP addresses to Ethernet addresses. The mapping is a cache, or “soft state”: the `NetworkInterface` keeps it around for efficiency’s sake, but if it has to restart from scratch, the mapping will naturally be regenerated without causing a problem.

```
void NetworkInterface::send_datagram(const InternetDatagram &dgram,
                                     const Address &next_hop)
```

This method is called when the caller (e.g., your `TCPConnection` or a router) wants to send an outbound Internet (IP) datagram to the next hop.¹It’s your interface’s job to translate this datagram into an Ethernet frame and (eventually) send it.

•*If the destination Ethernet address is already known*, send it right away. Create an Ethernet frame (with `type = EthernetHeader::TYPE_IPv4`), set the payload to be the serialized datagram, and set the source and destination addresses.

¹Please don’t confuse the ultimate the destination address, with the next hop. In this lab you’re going to care about the next hop’s address.

• *If the destination Ethernet address is unknown*, broadcast an ARP request for the next hop’s Ethernet address, and queue the IP datagram so it can be sent after the ARP reply is received.

Except: You don’t want to flood the network with ARP requests. If the network interface already sent an ARP request about the same IP address in the last five seconds, don’t send a second request—just wait for a reply to the first one. Again, queue the datagram until you learn the destination Ethernet address.

```
optional<InternetDatagram> NetworkInterface::recv_frame(const EthernetFrame &frame)
```

This method is called when an Ethernet frame arrives from the network. The code should ignore any frames not destined for the network interface (meaning, the Ethernet destination is either the

broadcast address or the interface's own Ethernet address stored in the `_ethernet_address` member variable).

- If the inbound frame is IPv4, parse the payload as an `InternetDatagram` and, if successful (meaning the `parse()` method returned `ParseResult::NoError`), return the resulting `InternetDatagram` to the caller.

- If the inbound frame is ARP, parse the payload as an `ARPMessage` and, if successful, remember the mapping between the sender's IP address and Ethernet address for 30 seconds. (Learn mappings from both requests and replies.) In addition, if it's an ARP request asking for our IP address, send an appropriate ARP reply.

```
void NetworkInterface::tick(const size_t ms_since_last_tick)
```

This is called as time passes. Expire any IP-to-Ethernet mappings that have expired.

You can test your implementation by running `ctest -V -R "^arp"`. This test does not rely on your TCP implementation.

4 webget revisited

Remember your `webget.cc` that you wrote in Lab 1 (using the Linux-provided TCP implementation in `TCPsocket`)? And, remember how you modified it in Lab 5 to use your own TCP-in-IP implementation in `CS144TCPsocket`? As we discussed above, this still relies on the Linux kernel for part of the stack: the network interface that translates between IP and the link layer (Ethernet).

We'd like you to switch it to use your network interface without changing anything else. All you'll need to do is replace the `CS144TCPsocket` type with `FullstackSocket`. This uses a TCP-in-IP-in-Ethernet stack as shown in Figure 1 (left side): your `webget.cc` application, on top of your `TCPConnection` implementation of TCP, on top of the TCP-in-IP code in `tcp_helpers/tcp_over_ip.cc`, on top of your `NetworkInterface`.

Recompile, and run `make check_lab6` to confirm that you've gone full-stack: you've written a basic Web fetcher on top of *your own complete TCP implementation and your own network interface implementation*, and it still successfully talks to a real webserver.

If you have trouble, try running the program manually:

```
./apps/webget server2.lab427.com /check_ans/ans , and try capturing what it's sending and receiving by using wireshark . You can save the packets it's sending and receiving by running  
sudo tcpdump -i tap10 -w /tmp/packets.tap . Then open the /tmp/packets.tap file in wireshark .
```

5 Q & A

•*How much code are you expecting?*

Overall, we expect the implementation (in `networkinterface.cc`) will require about 100–150 lines of code in total.

•*How does the `NetworkInterface` actually send an Ethernet frame?*

Similar story to how the `TCPSender` and `TCPConnection` worked. For the `NetworkInterface`, push any outbound frame to the `_frames_out` queue and it will be popped and sent by the object's owner.

•*What data structure should I use to record the mapping between next-hop IP address and Ethernet addresses?*

Up to you!

•*How do I convert an IP address that comes in the form of an `Address` object, into a raw 32-bit integer that I can write into the ARP message?*

Use the `Address::ipv4_numeric()` method.

•*What should I do if the `NetworkInterface` sends an ARP request but never gets a reply? Should I resend it after some timeout? Signal an error to the original sender using ICMP?*

In real life, yes, both of those things, but don't worry about that in this lab. (In real life, an interface will eventually send an ICMP "host unreachable" back across the Internet to the original sender if it can't get a reply to its ARP requests.)

•*What should I do if an Internet Datagram is queued waiting to learn the Ethernet address of the next hop, and that information never comes? Should I drop the datagram after some timeout?*

Again, definitely a "yes" in real life, but don't worry about that in this lab.

•*How do I run the test suite for this lab?*

`make check_lab6` (two tests). Or you can run the entire test suite with `make check` (160 tests).

6 Submit

1. In your submission, please only make changes to the `.hh` and `.cc` files in the top level of `libsponge`. Within these files, please feel free to add private members as necessary, but please don't change the *public* interface of any of the classes.
2. Before handing in any assignment, please run these in order:
 - a. `git status` (to check for un-committed changes - if you have any, commit!)
 - b. `make` (to make sure the code compiles)
 - c. `make check` (to make sure the automated tests pass)

3. Write a report in `writeups/lab6.md` . This file should be a roughly 20-to-50-line document with no more than 80 characters per line to make it easier to read. The report should contain the following sections:
 - a. **Program Structure and Design.** Describe the high-level structure and design choices embodied in your code. You do not need to discuss in detail what you inherited from the starter code. Use this as an opportunity to highlight important design aspects and provide greater detail on those areas for your grading TA to understand. You are strongly encouraged to make this writeup as readable as possible by using subheadings and outlines. Please do not simply translate your program into an paragraph of English.
 - b. **Implementation Challenges.** Describe the parts of code that you found most troublesome and explain why. Reflect on how you overcame those challenges and what helped you finally understand the concept that was giving you trouble. How did you attempt to ensure that your code maintained your assumptions, invariants, and preconditions, and in what ways did you find this easy or difficult? How did you debug and test your code?
 - c. **Remaining Bugs.** Point out and explain as best you can any bugs (or unhandled edge cases) that remain in the code.
4. Please also fill in the number of hours the assignment took you and any other comments.
5. When ready to submit, please follow the instructions in Our QQ group. Please make sure you have committed everything you intend before submitting.
6. Please also fill in the number of hours the assignment took you and any other comments. Good luck!